

Multi-Agent Systems

Homework Assignment 4

MSc AI, VU

E.J. Pauwels

Version: November 26, 2021 — Deadline: Friday, December 3, 2021 (23h59)

4 Monte Carlo simulation

4.1 MC sampling

Recall that Monte Carlo sampling allows us to estimate the expectation of a random function by sampling from the corresponding probability distribution. More precisely, if $f(x)$ is a 1-dim (continuous) probability density, and $X \sim f$ is a stochastic variable distributed according to this density f , then the expected value of some function φ can be estimated using Monte Carlo sampling by:

$$E_f(\varphi(X)) \equiv \int \varphi(x)f(x) dx \approx \frac{1}{n} \sum_{i=1}^n \varphi(X_i) \quad \text{for sample of independent } X_1, X_2, \dots, X_n \sim f.$$

1. Assume that $X \sim N(0,1)$ is standard normal. Estimate the mean value $E(\cos^2(X))$. Quantify the uncertainty on your result.
2. Suppose you're designing a deep neural network that needs to maximize some score function S . The actual design of the network depends on some hyperparameter A . Training the networks is computationally very demanding and time consuming, and as a consequence you have only been able to perform ten experiments to date. Based on these ten data points you observe a slight positive correlation of 0.3 between the value of the hyperparameter A and the score S . If this result is genuine, it suggest to increase A in the next experiment in order to improve the score. But if the correlation is not significant, increasing A could lead you astray. How would you use MC to decide whether the correlation is significant?

Hint: Compute the empirical p-value of the observed result, under the assumption of independence.

4.2 Importance Sampling

Importance sampling extends the basic MC approach to cases where it is difficult to sample from f but (relatively) easy to sample from a (somewhat) similar distribution g . More precisely:

$$\begin{aligned}
E_f(\varphi(X)) &= \int \varphi(x) f(x) dx \\
&= \int \varphi(x) \frac{f(x)}{g(x)} g(x) dx \equiv E_g \left[\varphi(X) \frac{f(X)}{g(X)} \right] \\
&\approx \frac{1}{n} \sum_{i=1}^n \varphi(X_i) \frac{f(X_i)}{g(X_i)} \quad \text{for sample of independent } X_1, \dots, X_n \sim g.
\end{aligned}$$

1. Let $X \sim N(0, 1)$ be a standard normal stochastic variable. Use importance sampling to estimate $E(X^2)$ by sampling from a uniform distribution $q \sim U(-5, 5)$ on the interval $[-5, 5]$. What value do you expect (based on your knowledge of the normal distribution)? How accurate is your estimate based on importance sampling?
2. Suppose some random process produces output $(-1 \leq X \leq 1)$ that is distributed according to the following continuous density:

$$f(x) = \frac{1 + \cos(\pi x)}{2} \quad (\text{for } -1 \leq x \leq 1).$$

Again we are interested in estimation $E(X^2)$. However, as this is not a standard distribution it makes sense to use importance sampling to estimate this value. Quantify the uncertainty on your result.

4.3 Kullback-Leibler divergence

The Kullback-Leibler (KL) divergence quantifies the similarity (or dissimilarity) of two probability densities. More specifically, given two continuous (1-dim) probability densities f, g , the KL-divergence is defined as:

$$KL(f||g) = \int_{-\infty}^{\infty} f(x) \log \left(\frac{f(x)}{g(x)} \right) dx \equiv E_f \left[\log \left(\frac{f(X)}{g(X)} \right) \right] \quad (1)$$

1. Let $f \sim N(\mu, \sigma^2)$ and $g \sim N(\nu, \tau^2)$ both be normal distributions. Express $KL(f||g)$ as a function of the means and variances of f and g . We mention in passing that the KL expression in eq.1 is called a **divergence** rather than a **distance** because it's not symmetric. Use the expression obtained above to convince yourself of this fact.
2. Check your theoretical result in (1) by computing a sample-based estimate of the KL-divergence (Monte Carlo simulation). Pick an appropriate sample size. Compare the MC estimate to the theoretical result.

5 Exploitation versus Exploration

5.1 UCB versus ϵ -greedy for k -bandit problem

Write a programme to experiment with the exploration/exploitation for the k -bandit problem (pick some value $5 \leq k \leq 20$). Assume that the arms (a) generate normally distributed rewards with

unit standard deviation, but different means $q(a)$ (e.g. randomly generated). Assume that in every single experiment the agent can take a total of $T = 1000$ actions (i.e. arm pulls). Let $L(t)$ be the expected total regret at time t , defined as:

$$\ell(t) = E \left(\sum_{i=1}^t (q^* - q(a_i)) \right)$$

- Compute the experimental $L(t)$ curves for different strategies (ϵ -greedy for different values of ϵ , UCB). Compare to the theoretical lower bound found by Lai-Robbins:

$$\ell(t) \geq A \log(t) \quad \text{where} \quad A = \sum_{a: \Delta_a \neq 0} \frac{\Delta_a}{KL(f_a || f_a^*)} \quad \text{and} \quad \Delta_a = q^* - q(a).$$

- Compute and compare the percentage correct decisions (selection of correct arm) under the different strategies (i.e. ϵ -greedy for different values of ϵ , UCB). What is the influence of the UCB hyper-parameter c ?

PS: *No need to submit code, only the results.*

SOLUTIONS

Preliminary remark: How to compute uncertainty on MC estimate?

Suppose we want to use Monte Carlo (MC) to estimate the expectation $E(\phi(X))$ where X has a known density function $f(x)$ and ϕ is a known function. To compute the MC estimate we draw a large sample X_1, X_2, \dots, X_n from the distribution f and compute the corresponding function values $\varphi_i = \phi(X_i)$ for $i = 1, \dots, n$. The mean of these function values is the MC estimate for the expectation:

$$E\phi(X) \approx \frac{1}{n}(\varphi_1 + \varphi_2 + \dots + \varphi_n).$$

To compute the uncertainty on the estimate in the RHS we notice that this RHS is a sample mean $\bar{\varphi}$ of the φ_i observations which means that variance is equal to the variance of the φ_i observations divided by sample size:

$$\text{Var}(\bar{\varphi}) = \frac{\text{Var}(\varphi_i)}{n}$$

or equivalently (in terms of the standard deviation):

$$\text{std}(\bar{\varphi}) = \frac{\text{std}(\varphi_i)}{\sqrt{n}} \quad \text{a.k.a. standard error (s.e.) on MC estimate.}$$

The variance (or standard deviation) in the RHS can be estimated by computing the variance (or standard deviation) of the sample $\varphi_1, \varphi_2, \dots, \varphi_n$. It's standard practice to use two or three times the standard error as a measure for the uncertainty on the MC estimate.

4.1 Mont Carlo Sampling

Estimate $E(\cos^2(X))$ for $X \sim N(0,1)$: Matlab code

```
sample_size = 10000;    % sample size for MC estimate

X = randn(sample_size,1); % random sample from N(0,1) population
F = cos(X).^2;          % compute function value at each sample point

% The MC estimate for m = E((cos(X))^2) is obtained by computing the
% sample average:

m_mc = mean(F);

% Since each sample point in F is and independent sample from cos(X).^2
% (where X ~ N(0,1), the standard deviation std(F) is an estimate of the
% corresponding population standard deviation. The corresponding standard
% deviation for the sample mean is therefore equal to std(F)/sqrt(sample_size)

m_mc_std = std(F)/sqrt(sample_size);
```

Conclusion Based on the above matlab code we conclude that $E(\cos^2(X)) \approx 0.5665$ (based on 10000 MC samples). Since the standard error (standard deviation for sample mean) equals 0.003 (approximately), we estimate the accuracy on the result as $3 \times 0.003 \approx 0.01$. Hence we conclude:

$$E(\cos^2(X)) \approx 0.57 \pm 0.01.$$

Correlation between score and hyperparameter

- Assume that there is no correlation;
- Use this assumption to draw random samples (of size 10) from this distribution and compute the correlation coefficient.
- Compare the observed result $r_{obs} = 0.3$ to the correlations for the simulated samples. Compute how “extreme” the observed result is (i.e. compute its p -value). If the p -value is small (e.g. $p < 0.05$) the observed trend is likely to be genuine.
- BONUS: Since we have no guarantee that the data points are distributed according to a normal distribution, one could try some additional simulations in which one uses other likely distributions, to investigate how this impacts on the conclusion. **However, for small samples most histograms would be compatible with the normal distribution.**

MATLAB code:

```
% We have 10 data points for which the observed correlation equals r_obs = 0.3.

r_obs = 0.3;
n = 10; % number of experimental data points

% Assume that there is no correlation between the two parameters, then the
% observed correlation is a random fluctuation. To test how likely this
% size of fluctuation is, we generate independent variables and tally how
% often a correlation of r_obs (or larger) is observed.

nr_samples = 1000;
Rho_MC = zeros(nr_samples,1);

for i = 1:nr_samples
    % Generated randomly distributed but independent samples for S and A
    S = randn(n,1);
    A = randn(n,1);
    % Compute and store the observed correlation coef for each sample
    Rho = corrcoef(A,S); % full correlaton matrix
    Rho_mc(i) = Rho(1,2) ; % correlation btw. variables 1 and 2 in corr matrix
end

% Compute the p-value of the observed value
```

Histogram of MC values for correlation coef (assuming independence)
p-value of observed correlation ($r_{obs}=0.3$) = 0.212




Figure 1: Histogram of MC-based correlation values

```
pval = length(find(Rho_mc > r_obs))/nr_samples;
```

Conclusion The MC simulation assumes that there is no correlation between the hyperparameter and the score. In that case it turns out that approximately 20% of the simulated correlation coefficients exceed the observed value $r_{obs} = 0.3$ (i.e. p -value equals 0.2). Hence, the observed correlation is not significant, and does therefore not provide proof for a positive trend.

4.2 Importance Sampling

Matlab code for estimating EX^2 using samples from uniform

```
% X ~ N(0,1), hence density = f(x) = 1/sqrt{2\pi} exp(-x^2/2)
% Density for uniform U(-5,5) : g(x) = 1/10;
%
% We need to estimate EX^2 = Var(X) = 1 by sampling from the uniform;
%
% This means that we need to sample say U ~ U(-5,5) and compute the sample
% value :
% F = phi(U) (f(U)/g(U)) where phi(u) = u^2
```

```

sample_size = 10000

U = 10*rand(sample_size,1)-5;

F = (U.^2) .* (10*normpdf(U));

mc_estimate = mean(F);
mc_population_std = std(F);
mc_estimate_std = std(F)/sqrt(sample_size)

```

Conclusion The matlab code above produces the following result (uncertainty = 3 s.e.):

$$\text{MC estimate} = 0.99 \pm 0.03$$

where the error margins are based on the fact that the standard error (standard deviation of the sample mean) equals $se = 0.0105$. We take three times this value to quantify the uncertainty. Notice that the theoretical value is given by $EX^2 = VarX = 1$ (since $EX = 0$).

Matlab code for estimating from unusual distribution

```

% Question 2:
%-----

% X is distributed according to density f(x) = (1+cos(pi*x))/2;
f = @(x) (1+cos(pi*x))/2; % definition of density

sample_size = 1000 % for MC sample

% Sample from uniform
%=====

U = 2*rand(sample_size,1)-1; % Uniform on -1, 1; density = 1/2

% Compute the function value (weighted with correction factor)

F = (U.^2) .* (f(U)/(1/2)); % compute the pointwise result;

% Compute mean, std and se.

mc_estimate = mean(F);
mc_population_std = std(F);
mc_estimate_std = std(F)/sqrt(sample_size);

% Sample from triangular
%=====

```

```

% Define triangular density  $g(x) = x+1$  (if  $x \leq 0$ ) and
%  $g(x) = 1-x$  (if  $x > 0$ ):

g = @(x) (x<=0).*(x+1) + (x>0).*(-x+1);    % density of the triangular density

% To create an observation from the triangular distribution,
% add two independent unifr
Z = rand(sample_size,1) + rand(sample_size,1) - 1;    %

% Compute function value (weighted with correction factor)
F = (Z.^2) .* (f(Z)./g(Z)); % compute the pointwise result;

% Compute mean, std and se
mc_estimate = mean(F);
mc_population_std = std(F);
mc_estimate_std = std(F)/sqrt(sample_size);

+++++
MATLAB code yields following result:
+++++

===== Sampling from uniform on [-1,1] =====
MC estimate = 0.13251 with s.e. = 0.0028724
===== Sampling from triangular density on [-1,1] =====
MC estimate = 0.13431 with s.e. = 0.0037221

```

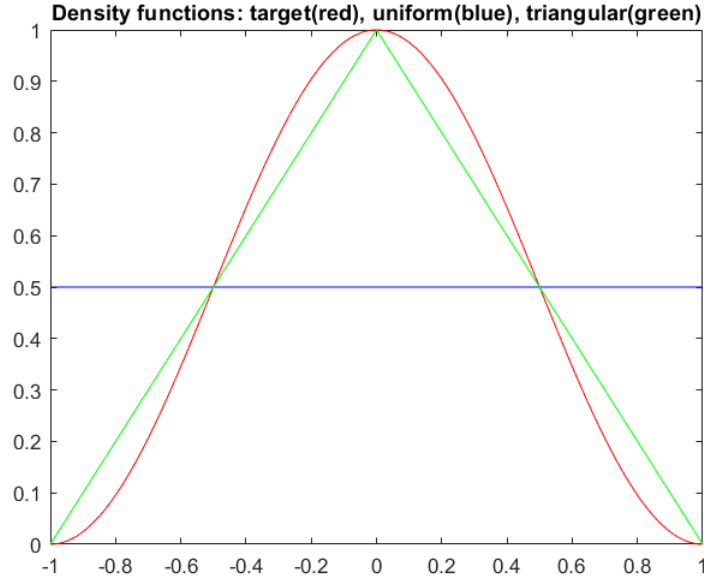



Figure 2: Target density (red) from which we need to sample. Two approximations used for importance sampling: uniform (blue) and triangular(green)

4.3 Kullback-Leibler divergence

KL for two gaussians Assuming normal densities $f \sim N(\mu_1, \sigma_1^2)$ and $g \sim N(\mu_2, \sigma_2^2)$, a straightforward computation yields:

$$\begin{aligned}
 KL(f||g) &= \int f(x) \log \left(\frac{f(x)}{g(x)} \right) dx \\
 \log \left(\frac{f(x)}{g(x)} \right) &= \log \left(\frac{e^{-(x-\mu_1)^2/2\sigma_1^2}}{\sqrt{2\pi}\sigma_1} \frac{\sqrt{2\pi}\sigma_2}{e^{-(x-\mu_2)^2/2\sigma_2^2}} \right) \\
 &= \log \left(\frac{\sigma_2}{\sigma_1} e^{-(x-\mu_1)^2/2\sigma_1^2 + (x-\mu_2)^2/2\sigma_2^2} \right) \\
 &= \log \left(\frac{\sigma_2}{\sigma_1} \right) + \log \left(e^{-(x-\mu_1)^2/2\sigma_1^2 + (x-\mu_2)^2/2\sigma_2^2} \right) \\
 &= \log \left(\frac{\sigma_2}{\sigma_1} \right) + \frac{-(x-\mu_1)^2}{2\sigma_1^2} + \frac{(x-\mu_2)^2}{2\sigma_2^2}
 \end{aligned}$$

$$\begin{aligned}
KL(f||g) &= \int f(x) \left(\log \left(\frac{\sigma_2}{\sigma_1} \right) - \frac{(x - \mu_1)^2}{2\sigma_1^2} + \frac{(x - \mu_2)^2}{2\sigma_2^2} \right) dx \\
&= \int f(x) \log \left(\frac{\sigma_2}{\sigma_1} \right) dx - \int f(x) \frac{(x - \mu_1)^2}{2\sigma_1^2} dx + \int f(x) \frac{(x - \mu_2)^2}{2\sigma_2^2} dx \\
&= \log \left(\frac{\sigma_2}{\sigma_1} \right) - \frac{\sigma_1^2}{2\sigma_1^2} + \int f(x) \frac{(x - \mu_1 + \mu_1 - \mu_2)^2}{2\sigma_2^2} dx \\
&= \log \left(\frac{\sigma_2}{\sigma_1} \right) - \frac{1}{2} + \int f(x) \frac{(x - \mu_1)^2 - 2(x - \mu_1)(\mu_1 - \mu_2) + (\mu_1 - \mu_2)^2}{2\sigma_2^2} dx \\
&= \log \left(\frac{\sigma_2}{\sigma_1} \right) - \frac{1}{2} + \frac{\sigma_1^2 - 0 + (\mu_1 - \mu_2)^2}{2\sigma_2^2} \\
&= \log \left(\frac{\sigma_2}{\sigma_1} \right) - \frac{1}{2} + \frac{\sigma_1^2 + (\mu_1 - \mu_2)^2}{2\sigma_2^2}
\end{aligned}$$

$$KL(f||g) = \int f(x) \log \left(\frac{f(x)}{g(x)} \right) dx = \log \frac{\sigma_2}{\sigma_1} + \frac{\sigma_1^2 + (\mu_1 - \mu_2)^2}{2\sigma_2^2} - \frac{1}{2}.$$

Notice the asymmetric role of both densities. Although not obvious from the above, the KL distribution is always non-negative.

MC for KL estimation

```

% f ~ N(mu1,sigma1^2)
mu1 = 0;    sigma1 = 2;

% g ~ N(mu2,sigma2^2)
mu2 = 2;    sigma2 = 3;

sample_size = 1000

% Sample from f and compute the KL value at each sample point

X = mu1 + sigma1*randn(sample_size,1);
KL = log(normpdf(X,mu1,sigma1)./normpdf(X,mu2,sigma2));

KL_div = mean(KL);    % KL divergence
KL_div_std = std(KL)/sqrt(sample_size);

% KL divergence based on theoretical expression:

KL_div_theory = log(sigma2/sigma1) + (sigma1^2+(mu1-mu2)^2)/(2*sigma2^2) - 1/2;

+++++ MATLAB OUTPUT +++++
Monte Carlo estimate of KL-divergence: 0.33517 with s.e.= 0.019843
Theoretical (exact) value = 0.34991
+++++

```

5.1 UCB versus ϵ -greedy for k -bandit problem

Lai-Robbins bound for two Gaussian For clarity's sake, we restrict our attention to two Gaussian densities f_1 and f_2 , both with unit variance, but different means. Let's denote the difference in the means as $\Delta = |\mu_1 - \mu_2|$. The KL divergence (see previous problem) is therefore given by:

$$KL(f_1||f_2) = 0 + \frac{1 + \Delta^2}{2} - \frac{1}{2} = \frac{\Delta^2}{2},$$

from which it follows that the Lai-Robbins coefficient is given by:

$$A = \frac{\Delta}{\Delta^2/2} = \frac{2}{\Delta}.$$

Question 1 Consider two Gaussians ($k = 2$) with randomly generated means -1.0430 and -0.3677 and unit variance. Hence, $\Delta = 0.6753$ and Lai-Robbins bound: $A_{LR} = 2.9615$. (see Fig 3). We use UCB hyperparameter $c = 2$

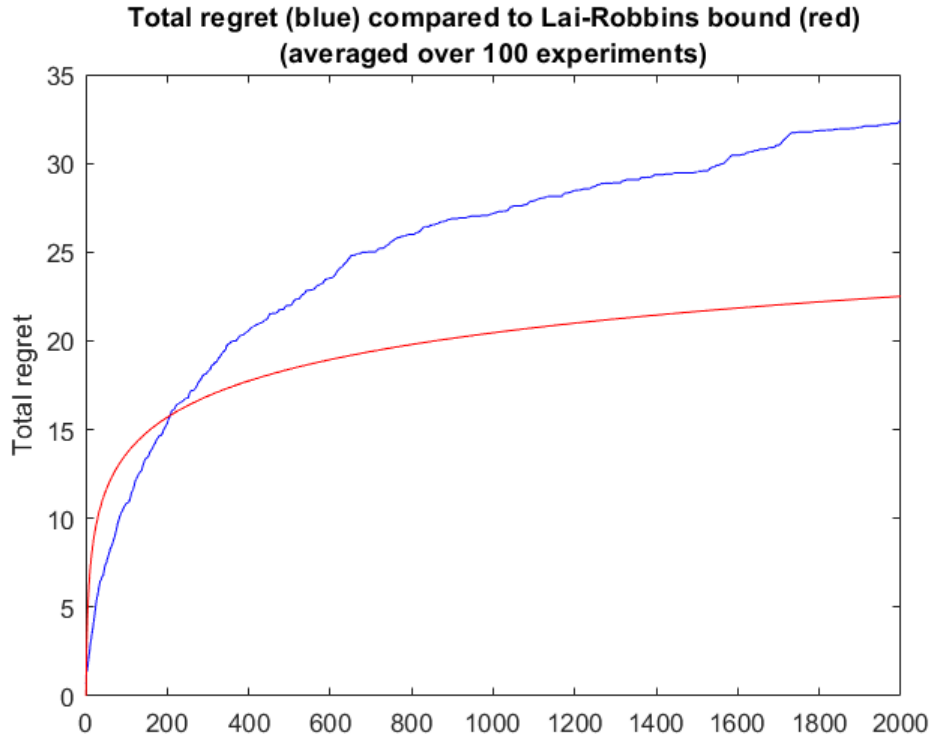


Figure 3: Total regret (blue) compared to Lai-Robbins (red)

Question 2 Comparing percentage of correct action choices for different algorithms (see Fig 4). Notice that by its very definition ϵ -greedy (with $\epsilon = 0.1$) cannot improve beyond the 90% level. The UCB performance clearly depends on the value of hyper-parameter c .

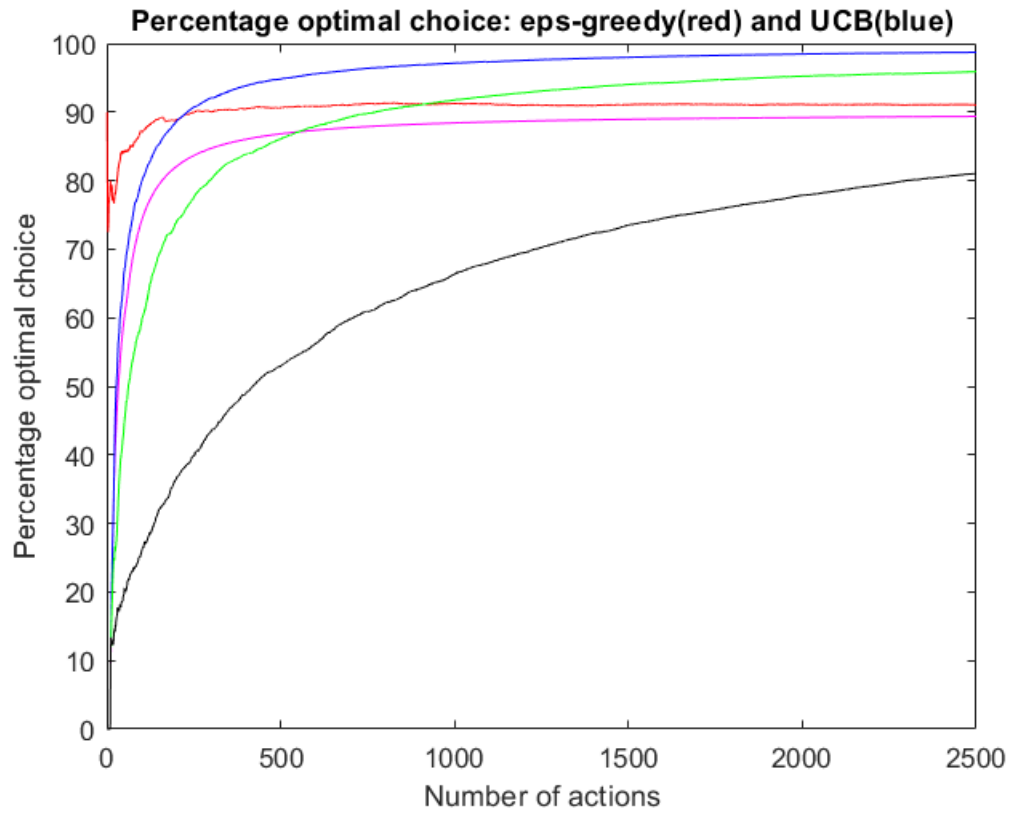


Figure 4: Comparing the percentage of correct choices for different exploration-exploitation strategies. The shown graphs are averaged over 10 experiments. Strategies: ϵ -greedy ($\epsilon = 0.1$): red, UCB with $c = 0.25$ (magenta), $c = 1$ (blue), $c = 2$ (green), $c = 5$ (black).