

Multi-Agent Systems

Introduction to Reinforcement Learning

Model-free Methods

Eric Pauwels (CWI & VU)

December 4, 2021

Outline

Policy Gradient Methods

Policy gradient algorithms

Policy gradient algorithms

- optimise policy directly,
- **NOT** via value function (indirectly)

Ingredients:

1. Parametrised policy $\pi_{\theta}(a|s)$ (θ to be determined)
2. Objective function $J(\theta)$ to be maximised
3. Update rule: $\theta_{new} \leftarrow \theta_{old}$, specifically **gradient** ascent:

$$\theta_{new} \leftarrow \theta_{old} + \nabla_{\theta} J(\theta_{old})$$

Policy gradient: Objective Function

- **Trajectory** (episodic):

$$\tau = \{s_0, a_0, r_1, s_1, a_1, r_2, s_2, \dots, s_{T-1}, a_{T-1}, r_T, s_T\}$$

- **Cumulative return along trajectory**

$$R(\tau) := \sum_{t=1}^T \gamma^{t-1} r_t \quad \text{or} \quad R_s(\tau) := \sum_{t=s}^T \gamma^{t-s} r_t$$

- **Objective function $J(\theta)$** : Quantifying policy performance:

$$J(\theta) := \mathbb{E}_{\tau \sim \pi_\theta} [R(\tau)]$$

- **Goal:**

$$\max_{\theta} J(\theta)$$

Policy gradient theorem

- Goal: maximise objective function

$$J(\theta) := \mathbb{E}_{\tau \sim \pi_{\theta}} [R(\tau)] = \int R(\tau) p(\tau | \theta) d\tau$$

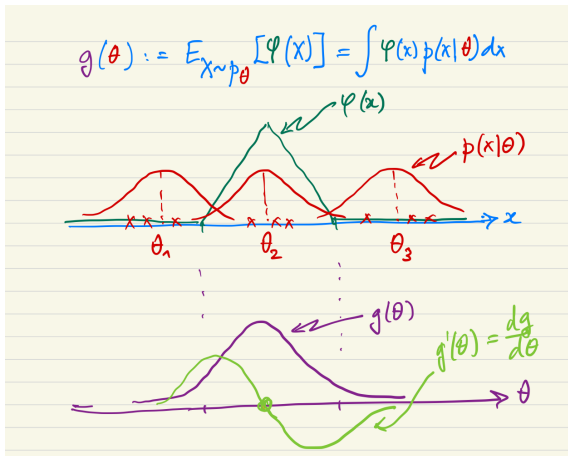
- In abstract terms (to simplify notation):

$$g(\theta) := \mathbb{E}_{x \sim p_{\theta}} [\phi(x)] = \int \phi(x) p(x | \theta) dx$$

- Need to compute derivative (to optimise):

$$\frac{d}{d\theta} g(\theta) = \frac{d}{d\theta} \int \phi(x) p(x | \theta) dx = \int \phi(x) \frac{d}{d\theta} (p(x | \theta)) dx$$

Policy gradient theorem



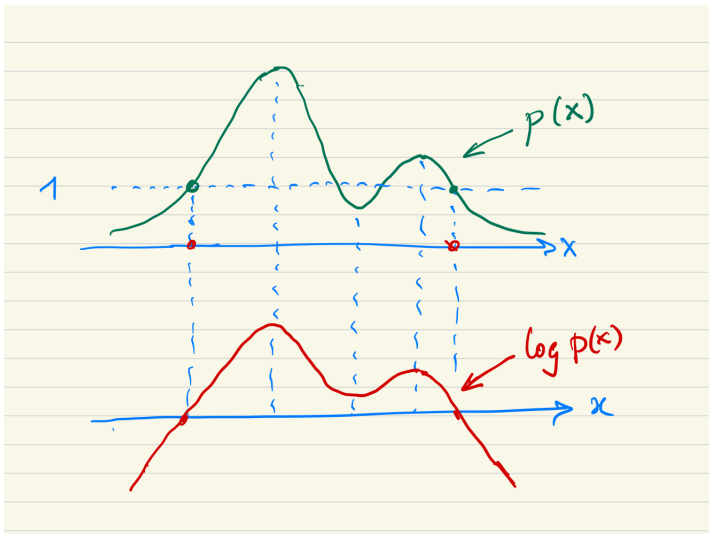
Optimal value $\theta^* = \theta_2$

Policy gradient theorem

$$\begin{aligned}
 \frac{d}{d\theta} g(\theta) &= \int \phi(x) \frac{d}{d\theta} (p(x|\theta)) dx \\
 &= \int \phi(x) \left[\frac{\frac{d}{d\theta} (p(x|\theta))}{p(x|\theta)} \right] p(x|\theta) dx \\
 &= \int \phi(x) \left[\frac{\frac{d}{d\theta} (p(x|\theta))}{p(x|\theta)} \right] p(x|\theta) dx \\
 &= \int \phi(x) \left[\frac{d}{d\theta} (\log p(x|\theta)) \right] p(x|\theta) dx \\
 &= \mathbb{E}_{X \sim p_\theta} \left[\phi(X) \frac{d}{d\theta} (\log p(x|\theta)) \right]
 \end{aligned}$$

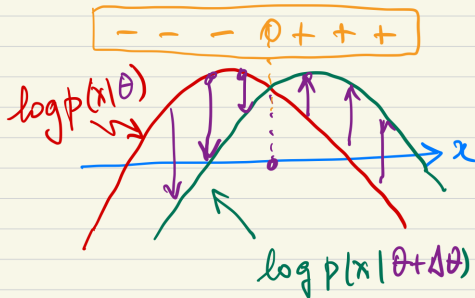
$$\boxed{\frac{d}{d\theta} E_{X \sim p_\theta} [\phi(X)] = \mathbb{E}_{X \sim p_\theta} \left[\phi(X) \frac{d}{d\theta} (\log p(X|\theta)) \right]}$$

$p(x)$ and $\log p(x)$ have same local extremes



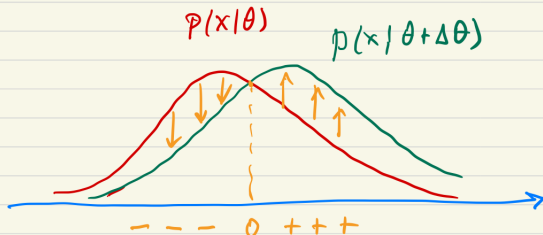
Policy gradient theorem

$$\nabla_{\theta} g(\theta) = \mathbb{E}_{x \sim \pi_{\theta}} [\varphi(x) \nabla_{\theta} \log p(x|\theta)]$$



at given x
 how does $\log p(x|\theta)$
 change if θ increases
 i.e. $\theta \rightarrow \theta + \Delta\theta$

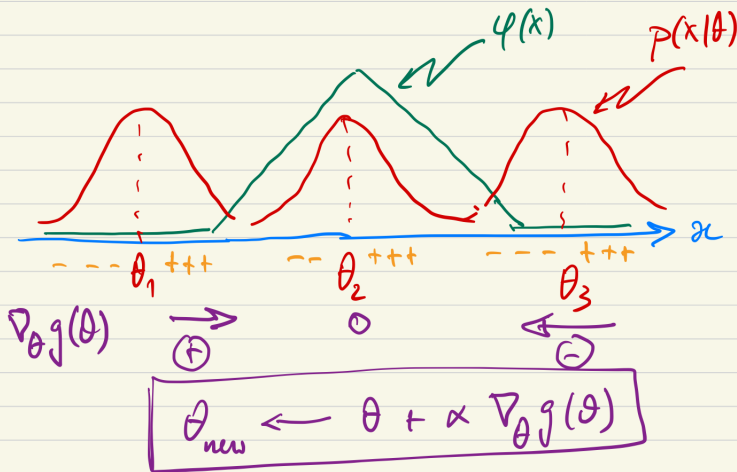
Policy gradient theorem



If high φ -values here
decrease θ
(neg. grad.)

If high φ -values here
increase θ
(pos. grad.)

Policy gradient theorem



Policy gradient theorem

General result

$$\nabla_{\theta} \mathbb{E}_{X \sim p_{\theta}} [\phi(X)] = \mathbb{E}_{X \sim p_{\theta}} [\phi(X) \nabla_{\theta} \log p(X | \theta)]$$

Gradient Policy Theorem

$$\nabla_{\theta} J(\theta) = \nabla_{\theta} \mathbb{E}_{\tau \sim \pi_{\theta}} [R(\tau)] = \mathbb{E}_{\tau \sim \pi_{\theta}} [R(\tau) \nabla_{\theta} \log p(\tau | \theta)]$$

Policy gradient theorem

$$p(\tau | \theta) = \prod_{t \geq 0} p(s_{t+1} | s_t, a_t) \pi_{\theta}(a_t | s_t)$$

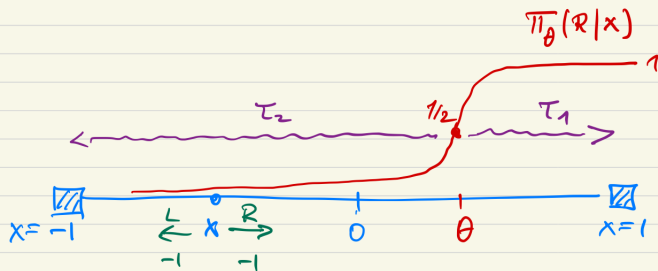
$$\log p(\tau | \theta) = \sum_{t \geq 0} [\log p(s_{t+1} | s_t, a_t) + \log \pi_{\theta}(a_t | s_t)]$$

$$\nabla_{\theta} \log p(\tau | \theta) = \sum_{t \geq 0} \nabla_{\theta} \log \pi_{\theta}(a_t | s_t)$$

Gradient Policy Theorem

$$\begin{aligned} \nabla_{\theta} J(\theta) &= \nabla_{\theta} \mathbb{E}_{\tau \sim \pi_{\theta}} [R(\tau)] = \mathbb{E}_{\tau \sim \pi_{\theta}} [R(\tau) \nabla_{\theta} \log p(\tau | \theta)] \\ &= \mathbb{E}_{\tau \sim \pi_{\theta}} \left[\sum_{t=0}^T R_t(\tau) \nabla_{\theta} \log \pi_{\theta}(a_t | s_t) \right] \end{aligned}$$

Policy gradient: Example(1)



$$R(\tau_1) = -(1-\theta)$$

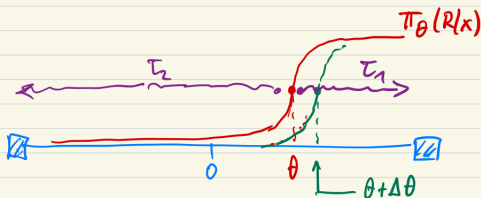
$$p(\tau_1|\theta) = 1/2$$

$$R(\tau_2) = -(1+\theta)$$

$$p(\tau_2|\theta) = 1/2$$

$$R(\tau_2) < R(\tau_1)$$

Policy gradient: Example(2)



$$p(\tau_1 | \theta + \Delta\theta) < p(\tau_1 | \theta) \rightarrow \nabla_{\theta} \log p(\tau_1 | \theta) < 0 \quad (= -\delta)$$

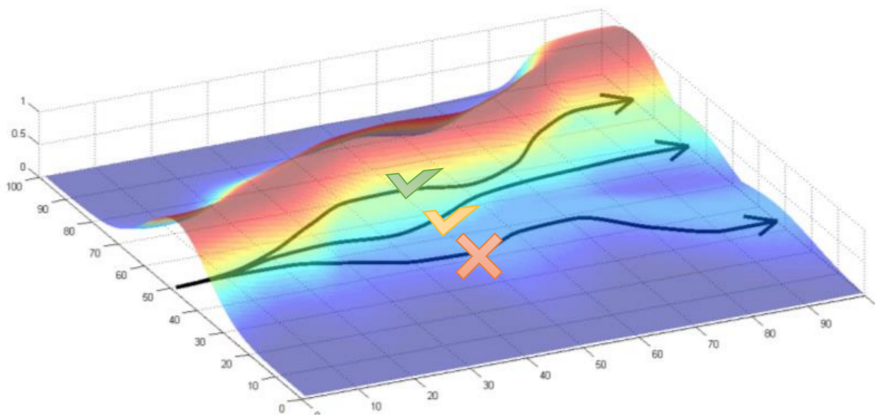
$$p(\tau_2 | \theta + \Delta\theta) > p(\tau_2 | \theta) \rightarrow \nabla_{\theta} \log p(\tau_2 | \theta) > 0 \quad (= +\delta)$$

$$\nabla_{\theta} J \approx R(\tau_1) \cdot \underbrace{\nabla_{\theta} \log p(\tau_1 | \theta)}_{-\delta} + R(\tau_2) \underbrace{\nabla_{\theta} \log p(\tau_2 | \theta)}_{+\delta}$$

$$\approx \delta [-R(\tau_1) + R(\tau_2)]$$

$$\approx \delta [|R(\tau_1)| - |R(\tau_2)|] < 0 \Rightarrow \theta \downarrow$$

Policy gradient illustration



REINFORCE algorithm

Example of policy gradient algo

1. Initialise learning rate α
2. Initialise parameter θ of policy π_θ
3. **for** episode = 1 ... NR_EPISODES:
4. Sample trajectory $\tau = \{s_0, a_0, r_1, s_1, \dots, r_T, s_T\}$
5. Set $\nabla_\theta J(\theta) = 0$
6. # add gradient contributions along trajectory
7. **for** $t = 0, 1, \dots, T$:
8. $R_t(\tau) = \sum_{s=t}^T \gamma^{t-s} r_s$,
9. $\nabla_\theta J(\theta) = \nabla_\theta J(\theta) + R_t(\tau) \nabla_\theta \log \pi_\theta(a_t | s_t)$
10. # update policy parameter
11. $\theta = \theta + \alpha \nabla_\theta J(\theta)$

Improving REINFORCE algorithm

- Policy gradient estimate has **high variance**
(trajectories might be very different!)

$$\nabla_{\theta} J(\theta) \approx \sum_{t=0}^T R_t(\tau) \nabla_{\theta} \log \pi_{\theta}(a_t | s_t)$$

- **Variance reduction** by introducing **action-independent baseline**:

$$\nabla_{\theta} J(\theta) \approx \sum_{t=0}^T (R_t(\tau) - b(s_t)) \nabla_{\theta} \log \pi_{\theta}(a_t | s_t)$$

- Example: **Actor-Critic algorithm**:

$$R_t(\tau) = q_{\pi_{\theta}}(s_t, a_t) \quad \text{and} \quad b(s_t) = v_{\pi_{\theta}}(s_t)$$

Gaussian Policy as example of parametrised policy

- In continuous action spaces, a Gaussian policy is natural
- Mean is a linear combination of state features $\mu(s) = \phi(s)^\top \theta$
- Variance may be fixed σ^2 , or can also be parametrised
- Policy is Gaussian, $a \sim \mathcal{N}(\mu(s), \sigma^2)$
- The score function is

$$\nabla_{\theta} \log \pi_{\theta}(s, a) = \frac{(a - \mu(s))\phi(s)}{\sigma^2}$$

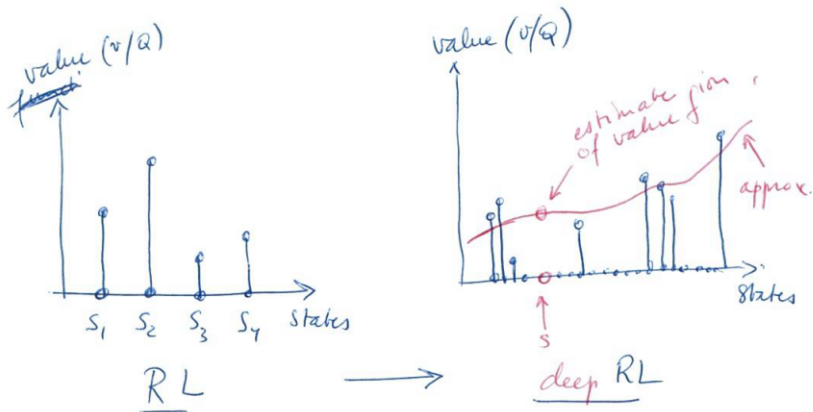
Outline

Policy Gradient Methods

Deep Q-Networks (DQN) as example of deep RL

Deep Reinforcement Learning (dRL)

- What if number of states or actions is **huge**!
- Use **function approximation** (i.e. **generalisation**) to estimate value functions in **unseen states**;



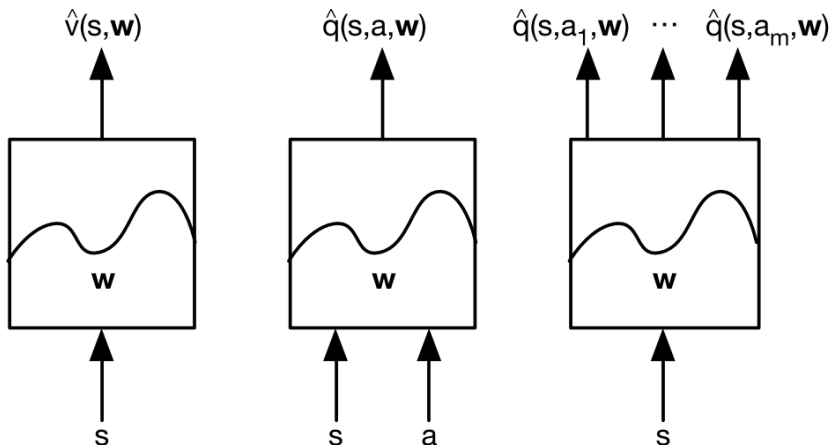
Value Function Approximation for Large RL Problems

- So far we have represented value function by a *lookup table*
 - Every state s has an entry $V(s)$
 - Or every state-action pair s, a has an entry $Q(s, a)$
- Problem with large MDPs:
 - There are too many states and/or actions to store in memory
 - It is too slow to learn the value of each state individually
- Solution for large MDPs:
 - Estimate value function with *function approximation*

$$\hat{v}(s, \mathbf{w}) \approx v_{\pi}(s)$$
$$\text{or } \hat{q}(s, a, \mathbf{w}) \approx q_{\pi}(s, a)$$

- *Generalise* from seen states to unseen states
- *Update* parameter \mathbf{w} using MC or TD learning

Types of Value Function Approximation



Fitting dNN using Stochastic Gradient Descent

- Goal: find parameter vector \mathbf{w} minimising mean-squared error between approximate value fn $\hat{v}(s, \mathbf{w})$ and true value fn $v_\pi(s)$

$$J(\mathbf{w}) = \mathbb{E}_\pi [(v_\pi(S) - \hat{v}(S, \mathbf{w}))^2]$$

- Gradient descent finds a local minimum

$$\begin{aligned}\Delta \mathbf{w} &= -\frac{1}{2} \alpha \nabla_{\mathbf{w}} J(\mathbf{w}) \\ &= \alpha \mathbb{E}_\pi [(v_\pi(S) - \hat{v}(S, \mathbf{w})) \nabla_{\mathbf{w}} \hat{v}(S, \mathbf{w})]\end{aligned}$$

- Stochastic gradient descent *samples* the gradient

$$\Delta \mathbf{w} = \alpha (v_\pi(S) - \hat{v}(S, \mathbf{w})) \nabla_{\mathbf{w}} \hat{v}(S, \mathbf{w})$$

- Expected update is equal to full gradient update

Incremental Prediction Algorithms

- Have assumed true value function $v_\pi(s)$ given by supervisor
- But in RL there is no supervisor, only rewards
- In practice, we substitute a *target* for $v_\pi(s)$
 - For MC, the target is the return G_t

$$\Delta \mathbf{w} = \alpha(G_t - \hat{v}(S_t, \mathbf{w})) \nabla_{\mathbf{w}} \hat{v}(S_t, \mathbf{w})$$

- For TD(0), the target is the TD target $R_{t+1} + \gamma \hat{v}(S_{t+1}, \mathbf{w})$

$$\Delta \mathbf{w} = \alpha(R_{t+1} + \gamma \hat{v}(S_{t+1}, \mathbf{w}) - \hat{v}(S_t, \mathbf{w})) \nabla_{\mathbf{w}} \hat{v}(S_t, \mathbf{w})$$

DQN Overview

- Deep RL version of Q-learning:
use **experience** to learn optimal $q^*(s, a)$
 - Construct dNN approximation \hat{q}_θ for q^*
 - Use experiences to train dNN \hat{q}_θ
- Off-policy algorithm
- Improve sample efficiency by storing experiences in experience replay buffer
- Target network improves stability

Experience Replay Memory (ERM)

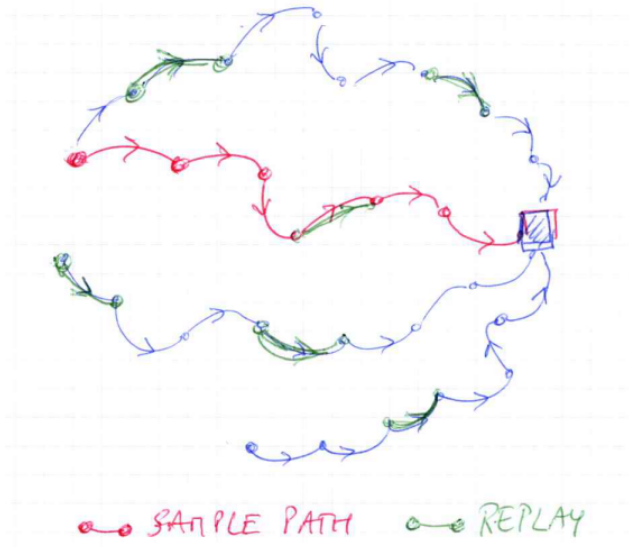
Motivation

- Recall: $\text{experience} = \{s, a, r, s'\}$
- Information-rich experiences should be used multiple times
- Experiences along trajectory are highly correlated;

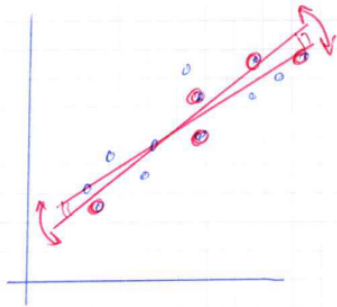
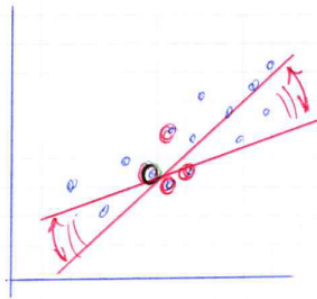
Experience Replay Memory (ERM)

- Storing: ERM stores K most recent experiences;
 - Typically: $K = 10^6$
 - Stalest data are overwritten (reflects learning)
- Training: sample batch (uniformly or prioritized) from ERM

Sample path versus experience replay



Correlated vs. uncorrelated



DQN Algorithm

1. **for** $m = 1 \dots \text{MAX_STEPS}$: *# loop over training steps;*
2. Use current ϵ -greedy policy to roll-out trajectories and store experiences $\{s, a, r, s'\}$ in ERM
3. **for** $b = 1 \dots B$: *# $B = \text{nr BATCHES}$ per training step*
4. *Sample batch of size U from ERM*
5. **for** $u = 1 \dots U$: *# $U = \text{number of updates per batch}$*
6. **for** $i = 1 \dots N$: *# $N = \text{batch size}$*
7. $y_i = r_i + \gamma \max_{a'_i} q_\theta(s'_i, a'_i)$ *# Target q -values*
8. $L(\theta) = \frac{1}{N} \sum_{i=1}^N (y_i - q_\theta(s_i, a_i))^2$ *# compute loss*
9. $\theta = \theta - \alpha \nabla_\theta L(\theta)$ *# Update dNN parameter θ*
10. Update ϵ, α

Improving DQN

- **Target Networks**

- Compute target values wrt. **lagged parameter value φ** of θ :

$$y_i = r_i + \gamma \max_{a'_i} q_{\varphi}(s'_i, a'_i)$$

- **Update $\varphi \leftarrow \theta$** from time to time;
- Fixes target, **improves stability**

- **PER: Prioritised Experience Replay**

- Not all experiences are equally informative to agent
- Therefore, **sampling uniformly** from ERM is **inefficient**
- Give **higher priority** to most “surprising” experiences (as measured by TD error).

DQN: Experience replay in Deep Q-Networks

- DQN uses experience replay and fixed Q-targets
- Store transition $(s_t, a_t, r_{t+1}, s_{t+1})$ in replay memory D
- Sample **random mini-batch** of transitions (s, a, r, s') from D
- Compute Q-learning targets w.r.t. old, fixed parameters w^-
- Optimize MSE between Q-network and Q-learning targets

$$\mathcal{L}_i(w_i) = \mathbb{E}_{s,a,r,s' \sim \mathcal{D}_i} \left[\left(\underbrace{r + \gamma \max_{a'} Q(s', a'; w_i^-)}_{\text{Q-learning target}} - \underbrace{Q(s, a; w_i)}_{\text{Q-network}} \right)^2 \right]$$

DQN: Deep Q-Network, (*Mnih et al, Nature 2015*)

- **Loss function**

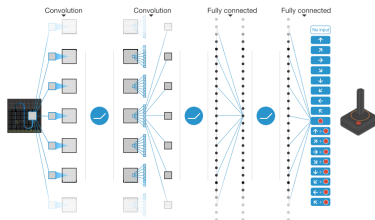
$$L_i(\theta_i) = \mathbb{E}_{(s,a,r,s') \sim U(D)} \left[\left(r + \gamma \max_{a'} Q(s', a'; \theta_i^-) - Q(s, a; \theta_i) \right)^2 \right]$$

- **Gradient**

$$\nabla_{\theta_i} L(\theta_i) = \mathbb{E}_{s,a,r,s'} \left[\left(r + \gamma \max_{a'} Q(s', a'; \theta_i^-) - Q(s, a; \theta_i) \right) \nabla_{\theta_i} Q(s, a; \theta_i) \right].$$

DQN: Deep Q-Network, (*Mnih et al, Nature 2015*)

- Single RL architecture achieved (super)human performance on suite of classic ATARI games;



- **Key features and improvements:**
 - Function approximation based on deep NN;
 - Experience replay (data re-use and de-correlation)
 - Use of target network for stabilisation