

Multi-Agent Systems

Introduction to Reinforcement Learning

Multi-Agent RL

Eric Pauwels (CWI & VU)

December 6, 2021

Outline

Actor-Critic

Advantage Actor-Critic (A2C)

Actor-Critic combines **valued-based** and **policy-based** learning.

Actor-Critic algorithms therefore have two components that are **learned jointly**:

- **Actor** learns a **parametrised policy**
- **Critic** learns **value function** to evaluate state-action pairs;

Advantage function: Select action based on how it **performs** **relative to other actions in that state**:

$$a_{\pi}(s, a) := q_{\pi}(s, a) - v_{\pi}(s)$$

Quick Recap

Policy gradient along trajectory τ

$$\nabla_{\theta} J(\theta) = \mathbb{E}_{\tau \sim \pi_{\theta}} \left[\sum_{t=0}^T R_t(\tau) \nabla_{\theta} \log \pi_{\theta}(a_t | s_t) \right]$$

Restricting the trajectory to the part starting at s_t :

$$R_t(\tau) = R(s_t, a_t, \dots, s_T) \implies \mathbb{E}_{\tau \sim \pi_{\theta}} [R_t(\tau)] = q_{\pi_{\theta}}(s_t, a_t);$$


$$\nabla_{\theta} J(\theta) = \mathbb{E}_{\tau \sim \pi_{\theta}} \left[\sum_{t=0}^T q_{\pi_{\theta}}(s_t, a_t) \nabla_{\theta} \log \pi_{\theta}(a_t | s_t) \right]$$

Policy Gradient

$$\nabla_{\theta} J(\theta) = E_{\tau \sim \pi_{\theta}} \left[\sum_{t=0}^T R_t(\tau) \nabla_{\theta} \log \pi_{\theta}(a_t | s_t) \right]$$

↓ single sample = single trajectory τ

$$\nabla_{\theta} J(\theta) \approx \sum_{t=0}^T R_t(\tau) \nabla_{\theta} \log \pi_{\theta}(a_t | s_t)$$



$$R_t(\tau) = r_{t+1} + r_{t+2} + \dots + r_T \approx q_{\pi_{\theta}}(s_t, a_t)$$

↑
generated
 $s_t \xrightarrow{a_t} s_{t+1}$

↑
1-sample.


Policy gradient: Quick Recap

$$\begin{aligned}\nabla_{\theta} J(\theta) &= \mathbb{E}_{\tau \sim \pi_{\theta}} \left[\sum_{t=0}^T R_t(\tau) \nabla_{\theta} \log \pi_{\theta}(a_t | s_t) \right] \\ &= \mathbb{E}_{\tau \sim \pi_{\theta}} \left[\sum_{t=0}^T q_{\pi_{\theta}}(s_t, a_t) \nabla_{\theta} \log \pi_{\theta}(a_t | s_t) \right]\end{aligned}$$

For N sampled paths $\tau_i = \{s_{i,0}, a_{i,0}, s_{i,1}, r_{i,1}, \dots\}$:

$$\begin{aligned}\nabla_{\theta} J(\theta) &\approx \frac{1}{N} \sum_{i=1}^N \left[\sum_{t=0}^T R_t(\tau_i) \nabla_{\theta} \log \pi_{\theta}(a_{i,t} | s_{i,t}) \right] \\ &= \frac{1}{N} \sum_{i=1}^N \left[\sum_{t=0}^T q_{\pi_{\theta}}(s_{i,t}, a_{i,t}) \nabla_{\theta} \log \pi_{\theta}(a_{i,t} | s_{i,t}) \right]\end{aligned}$$

Gradient policy theorem



Changing θ in the direction of $\nabla_\theta \log \pi(a|s)$ makes it more likely that action a will be chosen by policy π . That is a good thing if $q(s, a)$ positive/large!

Actor-Critic Methods

$$\nabla_{\theta} J(\theta) = \mathbb{E}_{\tau} \left[\sum_{t=0}^T \nabla_{\theta} \log \underbrace{\pi_{\theta}(a_t | s_t)}_{\text{ACTOR}} \underbrace{q_{\pi_{\theta}}(s_t, a_t)}_{\text{CRITIC}} \right]$$

- **Critic** estimates the **value function** (could be action-value q or state-value v function).
- **Actor** updates the **policy distribution** in the **direction suggested by the critic**.

Introducing baselines: A2C

- Policy gradient theorem:

$$\nabla_{\theta} J(\theta) = \mathbb{E}_{\tau} \left[\sum_{t=0}^T \nabla_{\theta} \log \pi_{\theta}(a_t | s_t) q(s_t, a_t) \right]$$

- **Problem:** value of $q(s, a)$ is not very informative;
- We need a reference point or **baseline**: **natural choice** = $v(s)$
- **Advantage:** Relative value of an action as compared to other actions in that state:

$$A(s, a) := q(s, a) - v(s)$$

- **Advantage actor-critic (A2C)**

$$\nabla_{\theta} J(\theta) \propto \mathbb{E}_{\tau} \left[\sum_{t=0}^T \nabla_{\theta} \log \pi_{\theta}(a_t | s_t) (q_{\pi_{\theta}}(s_t, a_t) - v_{\pi_{\theta}}(s_t)) \right]$$

Estimating Advantage

- Typically **two neural networks** to estimate
 1. **policy** $\rightarrow \pi_\theta$
 2. **value functions and advantage**: $\rightarrow v_w, q_w$
 3. Network weights: θ and w
- **Computational strategy**:
 - Estimate $v(s)$
 - Estimate $q(s, a)$ using Bellman eqs:

$$q(s_t, a_t) = \mathbb{E} [r_{t+1} + \gamma v(s_{t+1})]$$

- Along **sampled trajectory**:

$$\hat{q}(s_t, a_t) = r_{t+1} + \gamma \hat{v}(s_{t+1})$$

$$\hat{A}(s_t, a_t) = r_{t+1} + \gamma \hat{v}(s_{t+1}) - \hat{v}(s_t)$$


Estimating Advantage (2)

- n -step returns along sampled trajectory:

$$\hat{q}(s_t, a_t) = r_{t+1} + \gamma r_{t+2} + \gamma^2 r_{t+3} + \dots + \gamma^{n-1} r_{t+n} + \gamma^n \hat{v}(s_{t+n})$$

- Combination of biased and unbiased estimate:
 - Actual returns: **unbiased but high variance**
(sample paths can be very different!)
 - **Bias** due to inclusion of estimate \hat{v} , **but lower variance**
(average over all actions);

Mathematical aside (1): Bias vs. Variance



$$X_i \sim U(0, \theta)$$

↑ ?

$$\hat{\theta}_1 = \max_i X_i$$


BIASED
Small var

$$\hat{\theta}_2 = 2 \bar{X} = \frac{2}{n} \sum_{i=1}^n X_i$$

UNBIASED
Large VAR.

Mathematical aside (2a): Discount factor

Discount factor. $\gamma = \text{Prob of Continuing}$



$$\begin{aligned} ET &= \sum_{k=1}^{\infty} k \mathbb{P}(T=k) \\ &= \sum_{k=1}^{\infty} k \cdot \gamma^k (1-\gamma) \\ &= (1-\gamma) \gamma \underbrace{\sum_{k=1}^{\infty} k \gamma^{k-1}} \end{aligned}$$


Mathematical aside (2b): Discount factor

$$\begin{aligned}\sum_{k=1}^{\infty} k \gamma^{k-1} &= \frac{d}{d\gamma} \left(\sum_{k=1}^{\infty} \gamma^k \right) \\ &= \frac{d}{d\gamma} \left(\frac{\gamma}{1-\gamma} \right) \\ &= \frac{(1-\gamma) - \gamma(-1)}{(1-\gamma)^2} = \frac{1}{(1-\gamma)^2}\end{aligned}$$

$$ET = (1-\gamma) \gamma \frac{1}{(1-\gamma)^2} = \frac{\gamma}{1-\gamma}$$

Eg:
 $\gamma = 0.9$
 $ET \approx 9.$

Mathematical aside (2c): Discount factor



Further reading

[https://lilianweng.github.io/lil-log/2018/04/08/
policy-gradient-algorithms.html](https://lilianweng.github.io/lil-log/2018/04/08/policy-gradient-algorithms.html)